A Model-Based Variable Importance Metric for Supervised Learning

Brandon M. Greenwell*

Illumination Works

and

Andrew J. McCarthy

The Perduco Group

and

Bradley C. Boehmke

Department of Operational Sciences, Air Force Institute of Technology

October 5, 2017

Abstract

In the era of "big data", it is becoming more of a challenge to not only build state-of-the-art predictive models, but also gain an understanding of what's really going on in the data. For example, it is often of interest to know which, if any, of the predictors in a fitted model are relatively influential on the predicted outcome. Some modern algorithms—like random forests and boosted decision trees—have a natural way of quantifying the importance or relative influence of each feature. Other algorithms—like naive Bayes classifiers and support vector machines—are not capable of doing so and model-free approaches are generally used to measure each predictor's importance. In this paper, we discuss a model-based approach to measuring predictor importance in any supervised learning setting. The R code to reproduce all of the figures in this paper is available in the supplementary materials.

Keywords: Relative influence, Interaction, Partial dependence function, Partial dependence plot, PDP

1　Introduction

Complex supervised learning algorithms, such as neural networks and support vector machines, are more common than ever in predictive analytics, especially when dealing with large observational databases that don't adhere to the strict assumptions imposed by traditional statistical techniques (e.g., multiple linear regression which typically assumes linearity, homoscedasticity, and normality). However, it can be challenging to understand the results of such complex models and explain them to management. Graphical displays such as variable importance plots (when available) and partial dependence plots (PDPs) offer a simple solution.

PDPs are low-dimensional graphical renderings of the prediction function $\widehat{f}(\boldsymbol{x})$ so that the relationship between the outcome and predictors of interest can be more easily understood. These plots are especially useful in explaining the output from "black box" models. While PDPs can be constructed for any predictor in a fitted model, variable importance scores are more difficult to define. What does it mean for a feature to be important? It's influence on the predicted outcome? Some methods—like random forests and other tree-based methods—have a natural way of defining variable importance. In this paper, we consider a partial dependence-based variable importance measure that can be used with any supervised learning algorithm. Furthermore, we show how this new method can also be used to quantify the strength of potential interaction effects.

2　Background

We are often confronted with the task of extracting knowledge from large databases. For this task we turn to various statistical learning algorithms which, when tuned correctly, can have state-of-the-art predictive performance. However, having a model that predicts well is only solving part of the problem. It is also desirable to extract information about the relationships uncovered by the learning algorithm. For instance, we often want to know which predictors, if any, are important by assigning some type of variable importance score to each feature. Once a set of influential features has been identified, the next step is summarizing the functional relationship between each feature, or subset thereof, and the

outcome of interest. However, since most statistical learning algorithms are "black box" models, extracting this information is not always straightforward. Luckily, some learning algorithms have a natural way of defining variable importance.

2.1   Model-based approaches to variable importance

Decision trees probably offer the most natural model-based approach to quantifying the importance of each feature. In a binary decision tree, at each node t, a single predictor is used to partition the data into two homogeneous groups. The chosen predictor is the one that maximizes some measure of improvement $\widehat{i}_t$. The relative importance of predictor x is the sum of the squared improvements over all internal nodes of the tree for which x was chosen as the partitioning variable; see Breiman et al. (1984) for details. This idea also extends to ensembles of decision trees, such as boosting and random forest. In ensembles, the improvement score for each predictor is averaged across all the trees in the ensemble. Fortunately, due to the stabilizing effect of averaging, the improvement-based variable importance metric is often more reliable in large ensembles (Hastie et al., 2009, pg. 368). Random forests offer an additional method for computing variable importance scores. The idea is to use the leftover out-of-bag (OOB) data to construct validation-set errors for each tree. Then, each predictor is randomly shuffled in the OOB data and the error is computed again. The idea is that if variable x is important, then the validation error will go up when x is perturbed in the OOB data. The difference in the two errors is recorded for the OOB data then averaged across all trees in the forest.

In multiple linear regression, the absolute value of the t-statistic is commonly used as a measure of variable importance. The same idea also extends to generalized linear models and nonlinear least squares. Multivariate adaptive regression splines (MARS), which were introduced in Friedman (1991), is an automatic and adaptive regression technique which can be seen as a generalization of multiple linear regression and generalized linear models. In the MARS algorithm, the contribution (or variable importance score) for each predictor is determined using a generalized cross-validation (GCV) statistic.

For neural networks, two popular methods for constructing variable importance scores are the Garson algorithm (Garson, 1991), later modified by Goh (1995), and the Olden

algorithm (Olden et al., 2004). For both algorithms, the basis of these importance scores is the network's connection weights. The Garson algorithm determines variable importance by identifying all weighted connections between the nodes of interest. Olden's algorithm, on the other hand, uses the product of the raw connection weights between each input and output neuron and sums the product across all hidden neurons. This has been shown to outperform the Garson method in various simulations.

2.2 Filter-based approaches to variable importance

Filter-based approaches, which are described in Kuhn and Johnson (2013, chap. 18), do not make use of the fitted model to measure variable importance. They also do not take into account the other predictors in the model.

For regression problems, a popular approach to measuring the variable importance of a numeric predictor x is to first fit a flexible nonparametric model between x and the target Y; for example, the locally-weighted polynomial regression (LOWESS) method developed by Cleveland (1979). From this fit, a pseudo-R^2 measure can be obtained from the resulting residuals and used as a measure of variable importance. For categorical predictors, a different method based on standard statistical tests (e.g., t-tests and ANOVAs) is employed; see Kuhn and Johnson (2013, chap. 18) for details.

For classification problems, an area under the ROC curve (AUC) statistic can be used to quantify predictor importance. The AUC statistic is computed by using the predictor x as input to the ROC curve. If x can reasonably separate the classes of Y, then that is a clear indicator that x is an important predictor (in terms of class separation) and this is captured in the corresponding AUC statistic. For problems with more than two classes, extensions of the ROC curve or a one-vs-all approach can be used.

2.3 Partial dependence plots

Harrison and Rubinfeld (1978) analyzed a data set containing suburban Boston housing data from the 1970 census. They sought a housing value equation using an assortment of features; see Harrison and Rubinfeld (1978, Table IV) for a description of each variable. The usual regression assumptions, such as normality, linearity, and constant variance, were

clearly violated, but through an exhausting series of transformations, significance testing, and grid searches, they were able to build a model which fit the data reasonably well ($R^2 = 0.81$). Their prediction equation is given in Equation (1). This equation makes interpreting the model easier. For example, the average number of rooms per dwelling (RM) is included in the model as a quadratic term with a positive coefficient. This means that there is a monotonic increasing relationship between RM and the predicted median home value, but larger values of RM have a greater impact.

$$
\begin{aligned}
\log\widehat{(MV)} = {} & 9.76 + 0.0063RM^2 + 8.98 \times 10^{-5}AGE - 0.19\log\left(DIS\right) + 0.096\log\left(RAD\right) \\
& - 4.20 \times 10^{-4}TAX - 0.031PTRATIO + 0.36\left(B - 0.63\right)^2 - 0.37\log\left(LSTAT\right) \\
& - 0.012CRIM + 8.03 \times 10^{-5}ZN + 2.41 \times 10^{-4}INDUS + 0.088CHAS \\
& - 0.0064NOX^2.
\end{aligned}
\tag{1}
$$

Using modern computing power, many supervised learning algorithms can fit large data sets in seconds, producing powerful, highly accurate models. The downfall of many of these machine learning algorithms, however, is decreased interpretability. Unlike prediction formulas such as Equation (1), many machine learning algorithms are left attempting to convey some relative measure of variable importance which provide insight, but cannot match the simplicity of Equation (1) and others like it.

Quantifying the importance of each predictor is a crucial task in any supervised learning problem, but to gain even more insight, we can construct partial dependence plots (PDPs); see Friedman (2001) for details. PDPs are particularly effective at helping to explain the output from "black box" models, such as random forests and support vector machines. Not only do PDPs visually convey the relationship between low cardinality subsets of the feature set (usually 1-3) and the response (while accounting for the average effect of the other predictors in the model), they can also be used to rank and score the predictors in terms of their relative influence on the predicted outcome, as will be demonstrated in this paper.

Let $\boldsymbol{x} = \{x_1, x_2, \ldots, x_p\}$ represent the predictors in a model whose prediction function is $\widehat{f}(\boldsymbol{x})$. If we partition $\boldsymbol{x}$ into an interest set, $\boldsymbol{z}_s$, and its complement, $\boldsymbol{z}_c = \boldsymbol{x} \setminus \boldsymbol{z}_s$, then

the "partial dependence" of the response on z_s is defined as

$$f_s\left(\boldsymbol{z}_s\right) = E_{\boldsymbol{z}_c}\left[\widehat{f}\left(\boldsymbol{z}_s, \boldsymbol{z}_c\right)\right] = \int \widehat{f}\left(\boldsymbol{z}_s, \boldsymbol{z}_c\right) p_c\left(\boldsymbol{z}_c\right) d\boldsymbol{z}_c, \tag{2}$$

where $p_c\left(\boldsymbol{z}_c\right)$ is the marginal probability density of $\boldsymbol{z}_c$: $p_c\left(\boldsymbol{z}_c\right) = \int p\left(\boldsymbol{x}\right) d\boldsymbol{z}_s$. Equation (2) can be estimated from a set of training data by

$$\bar{f}_s\left(\boldsymbol{z}_s\right) = \frac{1}{n}\sum_{i=1}^{n}\widehat{f}\left(\boldsymbol{z}_s, \boldsymbol{z}_{i,c}\right), \tag{3}$$

where $\boldsymbol{z}_{i,c}$ $(i = 1, 2, \ldots, n)$ are the values of $\boldsymbol{z}_c$ that occur in the training sample; that is, we average out the effects of all the other predictors in the model.

Constructing a PDP (3) in practice is rather straightforward. To simplify, let $\boldsymbol{z}_s = x_1$ be the predictor variable of interest with unique values $\{x_{11}, x_{12}, \ldots, x_{1k}\}$. The partial dependence of the response on x_1 can be constructed as follows:

Input: the unique predictor values $x_{11}, x_{12}, \ldots, x_{1k}$;

Output: the estimated partial dependence values $\bar{f}_1\left(x_{11}\right), \bar{f}_1\left(x_{12}\right), \ldots, \bar{f}_1\left(x_{1k}\right)$.

for $i \in \{1, 2, \ldots, k\}$ **do**

(1) copy the training data and replace the original values of x_1 with the constant x_{1i};

(2) compute the vector of predicted values from the modified copy of the training data;

(3) compute the average prediction to obtain $\bar{f}_1\left(x_{1i}\right)$.

end

Algorithm 1: A simple algorithm for constructing the partial dependence of the response on a single predictor x_1.

The PDP for x_1 is obtained by plotting the pairs $\{x_{1i}, \bar{f}_1\left(x_{1i}\right)\}$ for $i = 1, 2, \ldots, k$. Algorithm (1) can be computationally expensive since it involves k passes over the training records. Some ideas are discussed in Greenwell (2017). Most importantly, Algorithm (1) is embarrassingly parallel and computing partial dependence functions for each predictor can be done rather quickly on a machine with a multi-core processor. For much larger data sets,

it may be worthwhile to reduce the grid size by using specific quantiles for each predictor, rather than using all the unique observations. For example, the partial dependence function can be approximated very quickly by using the min and max of each predictor along with the deciles of the unique predictor values. The exception is classification and regression trees based on single-variable splits which can make use of the efficient weighted tree traversal method described in Friedman (2001).

While PDPs are an invaluable tool in understanding the relationships uncovered by complex nonparametric models, they can be misleading in the presence of substantial interaction effects (Goldstein et al., 2015). To overcome this issue, Goldstein et al. developed the concept of individual conditional expectation (ICE) curves. ICE curves display the estimated relationship between the response and a predictor of interest for each observation; in other words, skipping step 1 (c) in Algorithm (1). Consequently, the PDP for a predictor of interest can be obtained by averaging the corresponding ICE curves across all observations. Although ICE curves provide a refinement over traditional PDPs in the presence of substantial interaction effects, in Section 3.2, we show how to use partial dependence functions to evaluate the strength of potential interactions effects.

2.4   Boston housing data

For illustration, we will use a corrected version of the Boston housing data analyzed in Harrison and Rubinfeld (1978); the data are available from Statlib at http://lib.stat.cmu.edu/datasets/boston_corrected.txt. Using the R package randomForest (Liaw and Wiener, 2002), we fit a random forest with tuning parameter $m_{try} = 6$ (chosen using 5-fold cross-validation) and 1000 trees. The model fit is reasonable, with an *out-of-bag* (pseudo) R^2 of 0.89. The random forest-based variable importance scores are displayed in Figure 1. Both plots indicate that the percentage of lower status of the population (lstat) and the average number of rooms per dwelling (rm) are highly associated with the median value of owner-occupied homes (cmedv). They also indicate that the proportion of residential land zoned for lots over 25,000 sq.ft (zn) has little association with cmedv.

The PDPs for these three variables are displayed in Figure 2. These plots indicate that both lstat and rm have a strong nonlinear relationship with the predicted outcome.
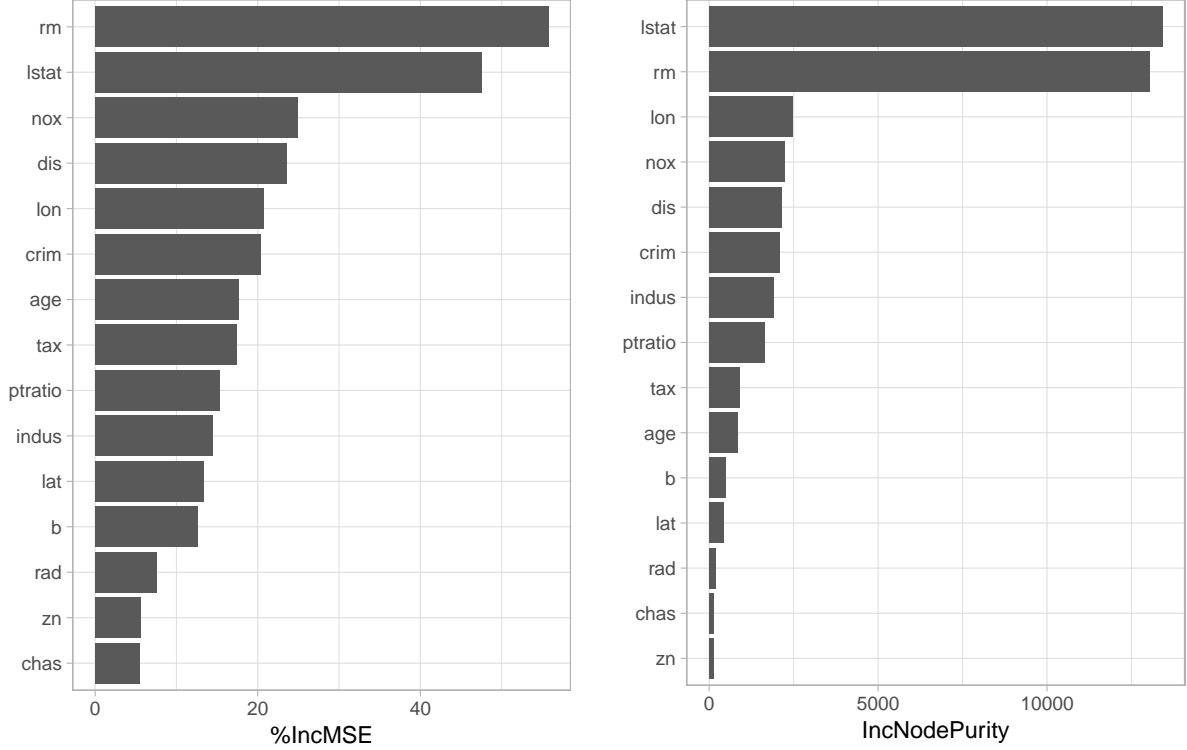
Figure 1: Variable importance scores from a random forest fit to the (corrected) Boston housing data. *Left*: OOB-based variable importance scores. *Right*: Impurity-based variable importance scores.

For instance, it seems that `rm` has relatively little effect on the predicted outcome until $rm \geq 5.5$, after which `rm` appears to have a monotonically increasing relationship with the predicted outcome. Notice how the PDP for `zn` is essentially flat, compared to the PDPs for `lstat` and `rm`. It is this notion of "flatness" which we will use as a basis to define our variable importance measure.

Figure 2: Partial dependence of `cmedv` on `lstat` (left), `rm` (middle), and `zn` (right). The mean of the $n = 506$ median home values is indicated by a dashed red line.

3 A partial dependence-based variable importance measure

The PDP for `zn` in Figure 2 is relatively flat (i.e., near-zero slope), indicating that `zn` does not have much influence on the predicted median home value. In other words, the partial dependence values $\bar{f}_i(x_{ij})$ $(j = 1, 2, \ldots, k_i)$ display little variability. One might conclude that any variable for which the PDP is "flat" is likely to be less important than those predictors whose PDP varies across a wider range of the response.

Our notion of variable importance is based on any measure of the "flatness" of the partial dependence function. In general, we define

$$\mathrm{imp}\,(x) = \mathrm{flatness}\left(\bar{f}_s\left(\boldsymbol{z}_s\right)\right),$$

where flatness is any measure of "flatness" of the curve. A simple and effective measure to use is the sample standard deviation for continuous predictors and the range statistic (divided by 4) for factors with K levels. Based on Algorithm (1), our importance measure for predictor x_1 is simply

$$\mathrm{imp}\,(x_1) = \begin{cases} \sqrt{\frac{1}{k-1} \sum_{i=1}^{k} \left[\bar{f}_1\left(x_{1i}\right) - \frac{1}{k}\sum_{i=1}^{k}\bar{f}_1\left(x_{1i}\right)\right]^2} & \text{if } x_1 \text{ is continuous} \\ \left[\max_i\left(\bar{f}_1\left(x_{1i}\right)\right) - \min_i\left(\bar{f}_1\left(x_{1i}\right)\right)\right]/4 & \text{if } x_1 \text{ is categorical} \end{cases}. \quad (4)$$

Note that our variable importance metric relies on the fitted model; hence, it is crucial to properly tune and train the model to have good performance.

To illustrate, we applied Algorithm (1) to all of the predictors in the random forest for the Boston housing example and compute (4). The results are displayed in Figure 3. In this case, our partial dependence-based algorithm matches closely with the results from the random forest. In particular, Figure 3 shows that `rm` and `lstat` are, by far, the most important variables in predicting median home value, while the predictors `zn`, `rad`, `lat`, and `b` seem to be the least important.


Figure 3: Partial-dependence-based variable importance scores from a random forest fit to the (corrected) Boston housing data.

3.1 Linear models

As mentioned earlier, a natural choice for measuring the importance of each term in a linear model is to use the absolute value of the corresponding coefficient divided by its estimated standard error (i.e., the absolute value of the t-statistic). This turns out to be equivalent to the partial dependence-based metric (4) when the predictors are independently and

uniformly distributed over the same range.

For example, suppose we have a linear model of the form

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon,$$

where β_i $(i = 1, 2)$ is a constant, X_1 and X_2 are both $\mathcal{U}(0, 1)$ random variables, and $\epsilon \sim \mathcal{N}(0, \sigma^2)$. Since we know the distribution of X_1 and X_2, we can easily find $f_1(X_1)$ and $f_2(X_2)$. For instance,

$$f_1(X_1) = \int_0^1 E[Y|X_1, X_2] p(X_2) dX_2,$$

where $p(X_2) = 1$. Simple calculus then leads to

$$f_1(X_1) = \beta_0 + \beta_2/2 + \beta_1 X_1 \quad and \quad f_2(X_2) = \beta_0 + \beta_1/2 + \beta_2 X_2..$$

Because $E[Y|X_1, X_2] = f(X_1, X_2)$ is additive, the true partial dependence functions are just simple linear regressions in each predictor with their original coefficient and an adjusted intercept. Taking the variance of each gives

$$Var[f_1(X_1)] = \beta_1^2/12 \quad and \quad Var[f_2(X_2)] = \beta_2^2/12.$$

Hence, the standard deviations are just the absolute values of the original coefficients (scaled by the same constant).

To illustrate, we simulated $n = 1000$ observations from the following linear model

$$Y = 1 + 3X_1 - 5X_2 + \epsilon,$$

where X_1 and X_2 are both $\mathcal{U}(0, 1)$ random variables, and $\epsilon \sim \mathcal{N}(0, 0.01^2)$. For this example, we have

$$f_1(X_1) = -\frac{3}{2} + 3X_1 \quad and \quad f_2(X_1) = \frac{5}{2} - 5X_2.$$

These are plotted as red lines in Figure 4. Additionally, the black lines in Figure 4 correspond to the estimated partial dependence functions using Algorithm (1).

Based on these plots, X_2 is clearly more influential than X_1. Taking the absolute value of the ratio of the slopes in $f_2(X_2)$ and $f_1(X_1)$ gives $5/3 \approx 1.67$. In other words, X_2 is roughly 1.67 times more influential on Y than X_1. Using the partial-dependence-based

11

Figure 4: Estimated and true partial dependence plots.

variable importance metric, we obtain $\operatorname{imp}(X_1) = 1.4828203$ and $\operatorname{imp}(X_2) = 0.8961719$ which gives the ratio $\operatorname{imp}(X_2) / \operatorname{imp}(X_1) \approx 1.65$. In fact, we can compute the true variance for both partial dependence functions:

$$Var\left[f_1\left(X_1\right)\right] = 3^2/12 \quad and \quad Var\left[f_2\left(X_2\right)\right] = 5^2/12.$$

Hence, the ratio of the true standard deviations is $5/3 \approx 1.67$.

Using the absolute value of the t-statistic becomes less useful in linear models when, for example, a predictor appears in multiple terms (e.g., interaction effects and polynomial terms). The partial dependence approach, on the other hand, does not suffer from such drawbacks.

3.2 Detecting interaction effects

As it turns out, our partial dependence-based variable importance measure (4) can also be used to quantify the strength of potential interaction effects. Let $\operatorname{imp}(x_i, x_j)$ $(i \neq j)$ be the standard deviation of the joint partial dependence values $\bar{f}_{ij}\left(x_{ii'}, x_{jj'}\right)$ for $i' = 1, 2, \ldots, k_i$ and $j' = 1, 2, \ldots, k_j$. Essentially, a weak interaction effect of x_i and x_j on Y would suggest that $\operatorname{imp}(x_i, x_j)$ has little variation when either x_i or x_j is held constant while the other

varies.

Let $z_s = (x_i, x_j)$, $i \neq j$, be any two predictors in the feature space $\boldsymbol{x}$. Construct the partial dependence function $\bar{f}_s(x_i, x_j)$ and compute $\text{imp}(x_i)$ for each unique value of x_j, denoted $\text{imp}(x_i|x_j)$, and take the standard deviation of the resulting importance scores. The same can be done for x_j and the results are averaged together. Large values (relative to each other) would be indicative of possible interaction effects.

4   Friedman's regression problem

To further illustrate, we will use one of the regression problems described in Friedman (1991) and Breiman (1996). The feature space consists of ten independent $\mathcal{U}(0, 1)$ random variables; however, only five out of these ten actually appear in the true model. The response is related to the features according to the formula

$$Y = 10 \sin(\pi x_1 x_2) + 20 (x_3 - 0.5)^2 + 10x_4 + 5x_5 + \epsilon,$$

where $\epsilon \sim \mathcal{N}(0, \sigma)$. Using the R package nnet (Venables and Ripley, 2002), we fit a neural network with one hidden layer containing eight units and a weight decay of 0.01 (these parameters were chosen using 5-fold cross-validation) to 500 observations simulated from the above model with $\sigma = 1$. The cross-validated R^2 value was 0.94.

Variable importance plots are displayed in Figure 5. Notice how the Garson and Olden algorithms incorrectly label some of the features not in the true model as "important". For example, the Garson algorithm incorrectly labels x_8 (which is not included in the true model) as more important than x_5 (which is in the true model). Similarly, Oden's method incorrectly labels x_{10} as being more important than x_2. Our method, on the other hand, clearly labels all five of the predictors in the true model as being, by far, the most important in the fitted neural network.

We also constructed the partial dependence functions for all pairwise interactions and computed the interaction statistic discussed in Section 3.2. The top ten interaction statistics are displayed in Figure 6. There is a clear indication of an interaction effect between the predictors x_1 and x_2, the only interaction present in the true model.

Figure 5: Variable importance plots for the neural network fit to the Friedman regression data. *Left*: partial dependence-based algorithm. *Middle*: Garson's algorithm. *Right*: Olden's algorithm.


Figure 6: Variable importance-based interaction statistics from the neural network fit to the Friedman regression data.

In fact, since we know the distributions of the predictors in the true model, we can work out the true partial dependence functions. For example, for the pairs (x_1, x_2) and (x_1, x_4), we have

$$f(x_1, x_2) = 10 \sin(\pi x_1 x_2) + 55/6,$$

14

and

$$f\left(x_1, x_4\right) = \frac{5\pi x_1\left(12x_4 + 5\right) - 12\cos\left(\pi x_1\right) + 12}{6\pi x_1}.$$

Next, we simulated the standard deviation of $f\left(x_1, x_2\right)$ for a wide range of fixed values of x_2; this is what $\text{imp}\left(x_1|x_2\right)$ is trying to estimate. The results from doing this for both predictors in each model are displayed in Figure 7. The top row of Figure 7 illustrates that the importance of x_1 (i.e., the strength of its relationship to the predicted outcome) heavily depends on the value of x_2 and vice versa (i.e., an interaction effect between x_1 and x_2). In the bottom row, on the other hand, we see that the importance of x_1 does not depend on the value of x_4 and vice versa (i.e., no interaction effect between x_1 and x_4).


Figure 7: Results from a small Monte Carlo simulation on the interaction effects between x_1 and x_2 (top row), and x_1 and x_4 (bottom row).

15

4.0.1 Friedman's H-statistic

An alternative measure for the strength of interaction effects is known as Friedman's H-statistic (Friedman and Popescu, 2008). Coincidentally, this method is also based on the estimated partial dependence functions of the corresponding predictors, but uses a different approach.

For comparison, we fit a generalized boosted regression model (GBM) to the Friedman regression data from the previous section. The parameters were chosen using 5-fold cross-validation. We used the R package `gbm` (Ridgeway, 2017) which has built-in support for computing Friedman's H-statistic for any combination of predictors. The results are displayed in Figure 8. To our surprise, the H-statistic did not seem to catch the true interaction between x_1 and x_2. Instead, the H-statistic ranked the pairs (x_8, x_9) and (x_7, x_{10}) as having the strongest interaction effects, even though these predictors do not appear in the true model. Our variable importance-based interaction statistic, on the other hand, clearly suggests the pair (x_1, x_2) as having the strongest interaction effect.

5 The Pima Indians diabetes data

In this section, we analyze the diabetes test results collected by the the US National Institute of Diabetes and Digestive and Kidney Diseases from a population of women who were at least 21 years old, of Pima Indian heritage, and living near Phoenix, Arizona. The data were taken directly from the `PimaIndiansDiabetes2` data frame which are available in the `mlbench` package (Leisch and Dimitriadou, 2012); for a description of the data and all the variables, see https://archive.ics.uci.edu/ml/datasets/pima+indians+diabetes.

We then used the `nnet` package (Venables and Ripley, 2002) in R to fit a neural network with seven hidden units and a weight decay of 0.01; these parameters were chosen to maximize AUC using repeated 5-fold cross-validation with ten repeats; the cross-validated AUC was 0.77.

The partial dependence functions for all the predictors are displayed in Figure 9. The PDPs clearly indicate that plasma glucose concentration (`glucose`), diabetes pedigree func-

Figure 8: Interaction statistics for the GBM model fit to the Friedman regression data. *Left*: Friedman's *H*-statistic. *Right*: Our variable importance-based interaction statistic.

tion (`pedigree`), and 2-hour serum insulin (`insulin`) have the largest effect on the neural network's prediction for the outcome of the diabetes test result (`diabetes`). On the other hand, the PDPs for diastolic blood pressure (`pressure`) and age (`age`) are relatively flat indicating they contribute less to the predictions than the other features. This is further illustrated in the variable importance plot displayed in Figure 10; the importance scores were obtained using Algorithm (4).

6 Conclusion

We have discussed a new partial dependence-based variable important measure for supervised learning. Since our algorithm is based on Friedman's idea of partial dependence

Figure 9: Partial-dependence plots for the eight predictors in the Pima Indians data set.


Figure 10: Partial-dependence-based variable importance scores for the eight predictors in the Pima Indians data set.

functions, our variable importance metric is model-based and takes into account the effects of the other features in the model. Consequently, this also requires that the fitted models be

properly tuned to achieve optimum performance. We have also discussed a simple extension of our variable importance metric that seems useful for detecting interaction effects.

While this new approach to appears to have high utility, more research is needed to determine where its deficiencies may lie. For example, outliers in the feature space can cause abnormally large fluctuations in the partial dependence values $\bar{f}(x_i)$ $(i = 1, 2, \ldots, k)$. Therefore, it may be advantageous to use more robust measures of spread to describe the variability in the estimated partial dependence values; a reasonable choice would be the median absolute deviation which has a finite sample breakdown point of $\lfloor n/2 \rfloor /n$. As mentioned in Greenwell (2017), it is also possible to replace the mean in step (3) of Algorithm (1) with a more robust estimate such as the median or trimmed mean. It would also be interesting to see if our metric could be extended to rely on individual conditional expectation (ICE) curves (Goldstein et al., 2015), a refinement over Friedman's PDP.

SUPPLEMENTARY MATERIAL

R package: The R package `vip`, hosted on GitHub at `https://github.com/AFIT-R/vip`, contains functions for computing variable importance scores and constructing variable importance plots for various types of fitted models in R using the method discussed in this paper.

R code: The R script 'greenwell-vi-2017.R' contains the R code to reproduce all of the results and figures in this paper.

References

Breiman, L., J. Friedman, and R. A. O. Charles J. Stone (1984). *Classification and Regression Trees*. The Wadsworth and Brooks-Cole statistics-probability series. Taylor & Francis.

Cleveland, W. S. (1979). Robust locally weighted regression and smoothing scatterplots. *Journal of the American Statistical Association 74* (368), 829–836.

Friedman, J. H. (1991). Multivariate adaptive regression splines. *The Annals of Statistics 19*(1), 1–67.

Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *The Annals of Statistics 29*, 1189–1232.

Friedman, J. H. and B. E. Popescu (2008). Predictive learning via rule ensembles. *Annals of Applied Statistics 2*(3), 916–954.

Garson, D. G. (1991). Interpreting neural-network connection weights. *Artificial Intelligence Expert 6*(4), 46–51.

Goh, A. (1995). Back-propagation neural networks for modeling complex systems. *Artificial Intelligence in Engineering 9*(3), 143–151.

Goldstein, A., A. Kapelner, J. Bleich, and E. Pitkin (2015). Peeking inside the black box: Visualizing statistical learning with plots of individual conditional expectation. *Journal of Computational and Graphical Statistics 24*(1), 44–65.

Greenwell, B. M. (2017). pdp: An r package for constructing partial dependence plots. *The R Journal 9*(1), 421–436.

Harrison, D. and D. L. Rubinfeld (1978). Hedonic housing prices and the demand for clean air. *Journal of Environmental Economics and Management 5*(1), 81–102.

Hastie, T., R. Tibshirani, and J. Friedman (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Second Edition*. Springer Series in Statistics. Springer-Verlag.

Kuhn, M. and K. Johnson (2013). *Applied Predictive Modeling*. SpringerLink : Bücher. Springer New York.

Leisch, F. and E. Dimitriadou (2012). *mlbench: Machine Learning Benchmark Problems*. R package version 2.1-1.

Liaw, A. and M. Wiener (2002). Classification and regression by randomforest. *R News 2*(3), 18–22.

Olden, J. D., M. K. Joy, and R. G. Death (2004). An accurate comparison of methods for quantifying variable importance in artificial neural networks using simulated data. *Ecological Modelling 178*(3), 389–397.

Ridgeway, G. (2017). *gbm: Generalized Boosted Regression Models*. R package version 2.1.3.

Venables, W. N. and B. D. Ripley (2002). *Modern Applied Statistics with S* (4th ed.). New York: Springer-Verlag.